



Bad news

We regret to inform you that our account on BuyMeACoffee has been suspended due to a violation of their terms of service. This was an unexpected development, and we are currently addressing the matter with utmost priority.

However, our mission at Freedium remains unchanged, and your support is more crucial than ever. We are transitioning to Patreon, a platform that aligns with our values and offers us the freedom to share our work with you.

Please join us on Patreon and continue to support our endeavors. Your contributions are invaluable to us, and we are committed to delivering the quality content you've come to expect from Freedium.

Thank you for your understanding and unwavering support.

Support Us on Patreon

Warm regards, The Freedium Team

[Patreon](#) - [Close](#) [Source code - Codeberg](#) - [Source code - GitHub](#)

[Freedium](#)



- [Source code](#)
- [Status page](#)
- [Patreon - Support us](#)

[< Go to the original](#)



Daily bit(e) of C++ | Optimizing code to run 87x faster

Daily bit(e) of C++ #474, Optimizing C++ code to run 87x faster for the One Billion Row Challenge (1brc).



[Šimon Tóth](#) [Follow](#)



[ITNEXT](#)

· ~16 min read · April 29, 2024 (Updated: April 30, 2024) · Free: Yes

The [One Billion Row Challenge](#) was initially a challenge for Java developers. The goal was to develop and optimize a parser for a file containing one billion records.

While the original challenge was for Java, this challenge is an excellent opportunity to demonstrate the optimization of C++ code and the related performance tools.

The challenge

Our input is a file called *measurements.txt*, which contains temperature measurements from various measurement stations. The file contains exactly one billion rows with the following format:

```
station name;value
station name;value
```

The station name is a UTF-8 string with a maximum length of 100 bytes, containing any 1-byte or 2-byte characters (except for " or '\n'). The measurement values are between -99.9 and 99.9, all with one decimal digit. The total number of unique stations is limited to 10⁵.

The output (to *stdout*) is a lexicographically sorted list of stations, each with the minimum, average and maximum measured temperature.

```
{Abha=-23.0/18.0/59.2, Abidjan=-16.2/26.0/67.3, Abéché=-10.0/29.4/69.0, ...}
```

Baseline implementation

Naturally, we have to start with a baseline implementation. Our program will have two phases: processing the input and formatting the output.

For the input, we parse the station name and the measured value and store them in a `std::unordered_map`.

```
struct Record {
    uint64_t cnt;
    double sum;

    float min;
    float max;
};

using DB = std::unordered_map<std::string, Record>;

DB process_input(std::istream &in) {
    DB db;

    std::string station;
    std::string value;

    // Grab the station and the measured value from the input
    while (std::getline(in, station, ';') &&
           std::getline(in, value, '\n')) {
        // Convert the measured value into a floating point
        float fp_value = std::stof(value);

        // Lookup the station in our database
        auto it = db.find(station);
        if (it == db.end()) {
            // If it's not there, insert
            db.emplace(station,
                       Record{1, fp_value, fp_value, fp_value});
            continue;
        }
        // Otherwise update the information
        it->second.min = std::min(it->second.min, fp_value);
        it->second.max = std::max(it->second.max, fp_value);
        it->second.sum += fp_value;
        ++it->second.cnt;
    }

    return db;
}
```

To produce the output, we collate the unique names, sort them lexicographically and then print out the minimum, average and maximum measurements.

```
void format_output(std::ostream &out, const DB &db) {
    std::vector<std::string> names(db.size());
    // Grab all the unique station names
    std::ranges::copy(db | std::views::keys, names.begin());
    // Sorting UTF-8 strings lexicographically is the same
    // as sorting by codepoint value
    std::ranges::sort(names);

    std::string delim = "";

    out << std::setiosflags(out.fixed | out.showpoint)
        << std::setprecision(1);

    out << "{";
    for (auto &k : names) {
        // Print StationName:min/avg/max
        auto &[_ , record] = *db.find(k);
        out << std::exchange(delim, ", ") << k << "="
            << record.min << "/" << (record.sum / record.cnt)
            << "/" << record.max;
    }
    out << "}\n";
}
```

We are left with a straightforward main function that plugs both parts together.

```
int main() {
    std::ifstream ifile("measurements.txt", ifile.in);
    if (not ifile.is_open())
        throw std::runtime_error("Failed to open the input file.");

    auto db = process_input(ifile);
    format_output(std::cout, db);
}
```

This baseline implementation is very simple, but sadly, it has two major problems: It's not correct (we will ignore that detail for now), and it's very, very slow. I will be presenting the performance on three machines:

- Intel 9700K on Fedora 39
- Intel 14900K on Windows Subsystem for Linux (Ubuntu)
- Mac Mini M1 (8-core)

Since we mainly care about the relative performance progression, the measurements are simply the fastest run on the corresponding machine.

- 9700K (Fedora 39): **132s**
- 14900K (WSL Ubuntu): **67s**
- Mac Mini M1: **113s**

Eliminating copies

We don't need special tooling for the first change. The most important rule for high-performance code is to avoid excessive copies. And we are making quite a few.

When we process a single line (and remember, there are 1 billion of them), we read the station name and the measured value into a `std::string`. This inherently introduces an explicit copy of the data because when we read from a file, the content of the file is already in memory in a buffer.

To eliminate this issue, we have a couple of options. We can switch to unbuffered reads, manually handle the buffer for the istream or take a more system-level approach, specifically memory mapping the file. For this article, we will go the memory-map route.

When we memory-map a file, the OS allocates address space for the file's content; however, data is only read into memory as required. The benefit is that we can treat the entire file as an array; the downside is that we lose control over which parts of the file are currently available in memory (relying on the OS to make the right decisions).

Since we are working with C++, let's wrap the low-level system logic in RAII objects.

```
// A move-only helper
template <typename T, T empty = T{}> struct MoveOnly {
    MoveOnly() : store_(empty) {}
    MoveOnly(T value) : store_(value) {}
    MoveOnly(MoveOnly &other)
        : store_(std::exchange(other.store_, empty)) {}
    MoveOnly &operator=(MoveOnly &other) {
        store_ = std::exchange(other.store_, empty);
        return *this;
    }
    operator T() const { return store_; }
    T get() const { return store_; }

private:
    T store_;
};

struct FileFD {
    FileFD(const std::filesystem::path &file_path)
        : fd_(open(file_path.c_str(), O_RDONLY)) {
        if (fd_ == -1)
            throw std::system_error(errno, std::system_category(),
                                     "Failed to open file");
    }

    ~FileFD() {
        if (fd_ >= 0)
            close(fd_);
    }

    int get() const { return fd_.get(); }

private:
    MoveOnly<int, -1> fd_;
};

struct MappedFile {
    MappedFile(const std::filesystem::path &file_path)
        : fd_(file_path) {
        // Determine the filesize (needed for mmap)
        struct stat sb;
        if (fstat(fd_.get(), &sb) == 1)
            throw std::system_error(errno, std::system_category(),
                                     "Failed to read file stats");

        sz_ = sb.st_size;

        begin_ = static_cast<char *>(
            mmap(NULL, sz_, PROT_READ, MAP_PRIVATE, fd_.get(), 0));
        if (begin_ == MAP_FAILED)
            throw std::system_error(errno, std::system_category(),
                                     "Failed to map file to memory");
    }

    ~MappedFile() {
        if (begin_ != nullptr)
            munmap(begin_, sz_);
    }

    // The entire file content as a std::span
    std::span<const char> data() const {
        return {begin_.get(), sz_.get()};
    }

private:
    FileFD fd_;
    MoveOnly<Char *> begin_;
    MoveOnly<size_t> sz_;
};
```

Because we are wrapping the memory-mapped file in a `std::span`, we can validate that everything still works by simply swapping `std::ifstream` for `std::ispanstream`.

```
int main() {
```

```

MappedFile mfile("measurements.txt");
std::istream ifile(mfile.data());

auto db = process_input(ifile);
format_output(std::cout, db);
}

```

While this validates that everything still works, it doesn't remove any excessive copies. To do that, we have to switch the input processing from operating on top of an istream to treating the input as one big C-style string.

```

DB process_input(std::span<const char> data) {
    DB db;

    auto iter = data.begin();

    while (iter != data.end()) {
        // Scan for the end of the station name
        auto semi_col = std::ranges::find(
            iter, std::unreachable_sentinel, ',');
        auto station = std::string_view(iter, semi_col);
        iter = semi_col + 1;

        // Scan for the end of measured value
        auto new_line = std::ranges::find(
            iter, std::unreachable_sentinel, '\n');
        // Parse as float
        float fp_value;
        std::from_chars(iter.base(), new_line.base(), fp_value);
        iter = new_line + 1;

        // Lookup the station in our database
        auto it = db.find(station);
        if (it == db.end()) {
            // If it's not there, insert
            db.emplace(station,
                Record{1, fp_value, fp_value, fp_value});
            continue;
        }
        // Otherwise update the information
        it->second.min = std::min(it->second.min, fp_value);
        it->second.max = std::max(it->second.max, fp_value);
        it->second.sum += fp_value;
        ++it->second.cnt;
    }

    return db;
}

```

We must adjust our hash map to support heterogeneous lookups because we now look up the stations using a `std::string_view`. This involves changing the comparator and adding a custom hasher.

```

struct StringHash {
    using is_transparent = void;

    size_t operator()(std::string_view txt) const {
        return std::hash<std::string_view>{}(txt);
    }
    size_t operator()(const std::string &txt) const {
        return std::hash<std::string>{}(txt);
    }
};

using DB = std::unordered_map<std::string, Record,
    StringHash, std::equal_to<>>;

```

This makes our solution run much faster (we are skipping over M1 for now since clang 15 doesn't support `std::from_chars` for floating point types).

- 9700K (Fedora 39): **47.6s**
- 14900K (WSL Ubuntu): **29.4s**

Analyzing the situation

To optimize the solution further, we have to analyze which parts of our solution are the main bottlenecks. We need a profiler.

When it comes to profilers, we have to choose between precision and low overhead. For this article, we will go with *perf*, a Linux profiler with extremely low overhead that still provides reasonable precision.

To have any chance to record a profile, we have to inject a bit of debugging information into our binary:

```

-fno-omit-frame-pointer # do not omit the frame pointer
-ggdb3 # detailed debugging information

```

To record a profile, we run the binary under the perf tool:

```
perf record --call-graph dwarf -F999 ./binary
```

The callgraph option will allow *perf* to attribute low-level functions to the correct caller using the debug information stored in the binary. The second option decreases the frequency at which perf captures samples; if the frequency is too high, some samples might be lost.

We can then view the profile:

```
perf report -g 'graph,caller'
```

However, if we run *perf* on the current implementation, we get a profile that isn't particularly informative.



We can deduce that the biggest portion of the runtime is spent in the `std::unordered_map`. However, the rest of the operations are lost in the low-level functions. For example, you might conclude that parsing out the measured values only takes 3% (the `std::from_chars` function); this would be an incorrect observation.

The profile is poor because we have put all the logic into a single tight loop. While this is good for performance, we completely lose the sense of the logical operations we are implementing:

- parse out the station name
- parse out the measured value
- store the data in the database

Profile clarity will drastically improve if we wrap these logical operations into separate functions.



Now, we can see that we are spending 62% of our time inserting data into our database, 26% parsing the measured values, and 5% parsing the station names.

We will address the hash map, but before that, let's deal with parsing the values. This will also fix a persistent bug in our code (bad rounding).

Integers in disguise

The input contains measurements from the range -99.9 to 99.9, always with one decimal digit. This means that we are not dealing with floating-point numbers in the first place; the measured values are fixed-point numbers.

The proper way to represent fixed-point values is as an integer, which we can manually parse (for now, in a straightforward way).

```
int16_t parse_value(std::span<const char>::iterator &iter) {
    bool negative = (*iter == '-');
    if (negative)
        ++iter;
    int16_t result = 0;
    while (*iter != '\n') {
        if (*iter != '.') {
            result *= 10;
            result += *iter - '0';
        }
        ++iter;
    }
    if (negative)
        result *= -1;
    ++iter;
    return result;
}
```

This change also propagates to the record structure.

```
struct Record {
    int64_t cnt;
    int64_t sum;

    int16_t min;
    int16_t max;
};
```

Database insertion can remain the same, but we can take the opportunity to optimize the code slightly.

```
void record(std::string_view station, int16_t value) {
    // Lookup the station in our database
    auto it = this->find(station);
    if (it == this->end()) {
        // If it's not there, insert
        this->emplace(station, Record{1, value, value, value});
        return;
    }
    // Switch minimum and maximum to exclusive branches
    if (value < it->second.min)
        it->second.min = value;
    else if (value > it->second.max)
        it->second.max = value;
    it->second.sum += value;
    ++it->second.cnt;
}
```

Finally, we have to modify the output formatting code. Since we are now working with fixed-point numbers, we have to correctly convert and round the stored integer values back to floating point.

```
void format_output(std::ostream &out, const DB &db) {
    std::vector<std::string> names(db.size());
    // Grab all the unique station names
    std::ranges::copy(db | std::views::keys, names.begin());
}
```

```

// Sorting UTF-8 strings lexicographically is the same
// as sorting by codepoint value
std::ranges::sort(names, std::less<>{});

std::string delim = "";

out << std::setiosflags(out.fixed | out.showpoint)
    << std::setprecision(1);
out << "{";
for (auto &k : names) {
    // Print StationName:min/avg/max
    auto &[_ , record] = *db.find(k);
    int64_t sum = record.sum;
    // CoRrect rounding
    if (sum > 0)
        sum += record.cnt / 2;
    else
        sum -= record.cnt / 2;

    out << std::exchange(delim, ", ") << k << "="
        << record.min / 10.0 << "/"
        << (sum / record.cnt) / 10.0 << "/"
        << record.max / 10.0;
}
out << "}\n";
}

```

This change fixes the aforementioned rounding bug and improves the runtime (floating-point arithmetic is slow). The implementation is also compatible with the M1 Mac.

- 9700K (Fedora 39): **35.5s (3.7x)**
- 14900K (WSL Ubuntu): **23.7s (2.8x)**
- Mac Mini M1: **55.7s (2.0x)**

Custom hash map

The `std::unordered_map` from the standard library is notorious for being slow. This is because it uses a node structure (effectively an array of linked lists of nodes). We could switch to a flat map (from Abseil or Boost). However, that would be against the original spirit of the 1brc challenge, which prohibited external libraries.

More importantly, our input is very constrained. We will have at most 10k unique keys for 1B records, leading to an exceptionally high hit ratio.

Because we are limited to 10k unique keys, we can use a linear probing hash map based on a 16-bit hash that directly indexes a statically sized array. We use the next available slot when we encounter a collision (two different stations map to the same hash/index).

This means that in the worst case (when all stations map to the same hash/index), we end up with a linear complexity lookup. However, that is exceptionally unlikely, and for the example input using `std::hash`, we end up with 5M collisions, i.e. 0.5%.

```

struct DB {
    DB() : keys_{}, values_{}, filled_{} {}

    void record(std::string_view station, int16_t value) {
        // Find the slot for this station
        size_t slot = lookup_slot(station);

        // If the slot is empty, we have a miss
        if (keys_[slot].empty()) {
            filled_.push_back(slot);
            keys_[slot] = station;
            values_[slot] = Record{1, value, value, value};
            return;
        }

        // Otherwise we have a hit
        if (value < values_[slot].min)
            values_[slot].min = value;
        else if (value > values_[slot].max)
            values_[slot].max = value;
        values_[slot].sum += value;
        ++values_[slot].cnt;
    }

    size_t lookup_slot(std::string_view station) const {
        // Get a hash of the name truncated to 16bit
        uint16_t slot = std::hash<std::string_view>{}(station);

        // While the slot is already occupied
        while (not keys_[slot].empty()) {
            // If it is the same name, we have a hit
            if (keys_[slot] == station)
                break;
            // Otherwise we have a collision
            ++slot;
        }

        // Either the first empty slot or a hit
        return slot;
    }

    // Helper method for the output formatting
    void sort_slots() {
        auto cmp = [this](size_t left, size_t right) {
            return keys_[left] < keys_[right];
        };
    }
};

```

```

    std::ranges::sort(filled_.begin(), filled_.end(), cmp);
}

// Keys
std::array<std::string, UINT16_MAX+1> keys_;
// Values
std::array<Record, UINT16_MAX+1> values_;
// Record of used indices (needed for output)
std::vector<size_t> filled_;
};

```

This change results in a sizeable speedup.

- 9700K (Fedora 39): **25.6s (5.1x)**
- 14900K (WSL Ubuntu): **18.4s (3.6x)**
- Mac Mini M1: **49.4s (2.3x)**

Micro-optimizations

We have cleared up the high-level avenues of optimizations, meaning it is time to dig deeper and micro-optimize the critical parts of our code.

Let's review our current situation.



We can potentially make some low-level optimizations in hashing (17%) and integer parsing (21%).

The correct tool for micro-optimizations is a benchmarking framework. We will implement several versions of the targetted function and compare the results against each other.

For this article, we will use Google Benchmark.

Parsing integers

The current version of integer parsing is (deliberately) poorly written and has excessive branching.

We can't properly make use of wide instructions (AVX) since the values can be as short as three characters. With wide instructions out of the way, the only approach that eliminates branches is a lookup table.

As we parse the number, we have only two possible situations (ignoring the sign):

- we encounter a digit: multiply the accumulator by 10 and add the digit value
- we encounter a non-digit: multiply the accumulator by 1 and add 0

We can encode this as a 2D compile-time generated array that contains information for all 256 values of a *char* type.

```

constexpr auto int_parse_table() {
    std::array<std::array<int16_t, 2>, 256> data;
    for (size_t c = 0; c < 256; ++c) {
        if (c >= '0' && c <= '9') {
            data[c][0] = c - '0';
            data[c][1] = 10;
        } else {
            data[c][0] = 0;
            data[c][1] = 1;
        }
    }
    return data;
}

static constexpr auto params = int_parse_table();

int16_t parse_int_table(const char *iter) {
    char sign = *iter;
    int16_t result = 0;
    while (*iter != '\n') {
        result *= params[*iter][1];
        result += params[*iter][0];
        ++iter;
    }
    ++iter;
    if (sign == '-')
        return result * -1;
    return result;
}

```

We can plug these two versions into our microbenchmark and get a very decisive result.



Sadly, the previous sentence is a lie. You cannot just plug the two versions into Google Benchmark. Our implementation is a tight loop over (at most) 5 characters, which makes it incredibly layout-sensitive. We can align the functions using an LLVM flag.

```
-mllvm -align-all-functions=5
```

However, even with that, the results fluctuate heavily (up to 40%).

Hashing

We have two opportunities for optimization when it comes to hashing.

Currently, we first parse out the name of the station, and then later, inside `lookup_slot`, we compute the hash. This means we traverse the data twice.

Additionally, we compute a 64-bit hash but only need a 16-bit one.

To mitigate the issues we run into with integer parsing, we will merge the parsing into a single step, producing a *string_view* of the station name, a 16-bit hash and the fixed-point measured value.

```
struct Measurement {
    std::string_view name;
    uint16_t hash;
    int16_t value;
};

Measurement parse_v1(std::span<const char>::iterator &iter) {
    Measurement result;

    auto end = std::ranges::find(iter, std::unreachable_sentinel, ';');
    result.name = {iter.base(), end.base()};
    result.hash = std::hash<std::string_view>{}(result.name);
    iter = end + 1;

    result.value = parse_int_table(iter);

    return result;
}
```

We use a simple formula to calculate the custom 16-bit hash and rely on unsigned overflow instead of modulo.

```
struct Measurement {
    std::string_view name;
    uint16_t hash;
    int16_t value;
};

Measurement parse_v2(std::span<const char>::iterator &iter) {
    Measurement result;

    const char *begin = iter.base();
    result.hash = 0;
    while (*iter != ';') {
        result.hash = result.hash * 7 + *iter;
        ++iter;
    }
    result.name = {begin, iter.base()};
    ++iter;

    result.value = parse_int_table(iter);

    return result;
}
```

This provides a nice speedup (with reasonable stability).

And when we plug this improvement into our solution, we get an overall speedup.

- 9700K (Fedora 39): **19.2s (6.87x)**
- 14900K (WSL Ubuntu): **14.1s (4.75x) (noisy)**
- Mac Mini M1: **46.2s (2.44x)**

Unleash the threads

If we investigate the profile now, we will see that we are reaching the limit of what is feasible.

There is very little runtime left outside of the parsing (which we just optimized), slot lookup and data insertion. So naturally, the next step is to parallelize our code.

The simplest way to achieve this is to chunk the input into roughly identical chunks, process each chunk in a separate thread and then merge the result.

We can extend our *MappedFile* type to provide chunked access.

```
struct MappedFile {
    /* ... */

    // Split the input into chunks
}
```



```

std::vector<std::span<const char>> chunked(size_t chunks) const {
    std::vector<std::span<const char>> result;

    size_t chunk_sz = sz_ / chunks;
    const char *chunk_begin = begin_;
    for (size_t i = 0; i < chunks - 1; ++i) {
        auto end = chunk_begin + chunk_sz;
        // align the end of the chunk to '\n' character
        while (end != begin_ + sz_ && *end != '\n')
            ++end;
        ++end;
        result.push_back({chunk_begin, end});
        chunk_begin = end;
    }
    // Manually set the last chunk
    result.push_back({chunk_begin, begin_ + sz_});

    return result;
}

private:
    MoveOnly<char *> begin_;
    MoveOnly<size_t> sz_;
};

```

We can then simply run our existing code per chunk, each in its own thread.

```

std::unordered_map<std::string, Record>
process_parallel(std::vector<std::span<const char>> chunks) {
    // Process the chunks in separate thread each
    std::vector<std::jthread> runners(chunks.size());
    std::vector<DB> dbs(chunks.size());
    for (size_t i = 0; i < chunks.size(); ++i) {
        runners[i] = std::jthread(
            [&, idx = i]() { dbs[idx] = process_input(chunks[idx]); });
    }
    runners.clear(); // join threads

    // Merge the partial DBs
    std::unordered_map<std::string, Record> merged;
    for (auto &db_chunk : dbs) {
        for (auto idx : db_chunk.filled_) {
            auto it = merged.find(db_chunk.keys_[idx]);
            if (it == merged.end()) {
                merged.insert_or_assign(db_chunk.keys_[idx],
                                         db_chunk.values_[idx]);
            } else {
                it->second.cnt += db_chunk.values_[idx].cnt;
                it->second.sum += db_chunk.values_[idx].sum;
                it->second.max = std::max(it->second.max,
                                         db_chunk.values_[idx].max);
                it->second.min = std::min(it->second.min,
                                         db_chunk.values_[idx].min);
            }
        }
    }
    return merged;
}

```

This gives a fairly nice scaling.



The following are the best results. Note that these are just best runs for relative comparison, not a rigorous benchmark.

- 9700K (Fedora 39): **2.6s (50x) (on 8 threads)**
- 14900K (WSL Ubuntu): **0.89s (75x) (on 32 threads)**
- Mac Mini M1: **10.2s (11x) (on 24 threads)**

Dealing with asymmetric processing speeds

The 9700K scales extremely cleanly, but that is because this processor has 8 identical cores that do not support hyperthreading. Once we move on to 14900K, the architecture gets a lot more complicated with performance and efficiency cores.

If we trivially split the input into identical chunks, the efficiency cores will trail behind and slow down the overall runtime. So, instead of splitting the input into chunks, one for each thread, let's have the threads request chunks as needed.

```

// Process the chunks in separate thread each
std::vector<std::jthread> runners(chunks);
std::vector<DB> dbs(chunks);
for (size_t i = 0; i < chunks; ++i) {
    runners[i] = std::jthread([&, idx = i]() {
        auto chunk = file.next_chunk();
        while (not chunk.empty()) {
            process_input(dbs[idx], chunk);
            chunk = file.next_chunk();
        }
    });
}
runners.clear(); // join threads

```

And the corresponding *next_chunk* method in our *MappedFile*.

```

std::span<const char> next_chunk() {
    std::lock_guard lock(mux_);
    if (chunk_begin_ == begin_ + sz_)
        return {};
}

```

```

size_t chunk_sz = 64 * 1024 * 1024; // 64MB

const char *end = nullptr;
// prevent reading past the end of the file
if (chunk_begin_ + chunk_sz > begin_ + sz_) {
    end = begin_ + sz_;
} else {
    end = chunk_begin_ + chunk_sz;
    while (end != begin_ + sz_ && *end != '\n')
        ++end;
    ++end;
}
std::span<const char> result(chunk_begin_, end);
chunk_begin_ = end;
return result;
}

```

Which allows us to squeeze the last bit of performance from the 14900K.

- 14900K (WSL Ubuntu): **0.77s (87x) (on 32 threads)**

Conclusion

We have improved the performance of our initial implementation by 87x, which is definitely not insignificant. Was it worth it?

Well, that depends. This article has taken me a long time to write and has cost me a chunk of my sanity. Notably, the alignment issues when using microbenchmarks were a huge pain to overcome.

If I were optimizing a production piece of code, I would likely stop at the basic optimizations (and threads). Micro-optimizations can be worthwhile, but the time investment is steep, and the stability of these optimizations on modern architectures is very poor.

The full source code is available in [this GitHub repository](#).

[#cpp#cplusplus#programming#software-development#performance](#)

Reporting a Problem

Sometimes we have problems displaying some Medium posts.

If you have a problem that some images aren't loading - try using VPN. Probably you have problem with access to Medium CDN (or fucking Cloudflare's bot detection algorithms are blocking you).

Describe your problem

Problem Description

Submit

Cancel